

# Data Wrangling

Estimated time needed: **30** minutes

## Objectives

After completing this lab you will be able to:

- Handle missing values
- Correct data format
- Standardize and Normalize Data

## Table of content

- [Identify and handle missing values](#)
  - [Identify missing values](#)
  - [Deal with missing values](#)
  - [Correct data format](#)
- [Data standardization](#)
- [Data Normalization \(centering/scaling\)](#)
- [Binning](#)
- [Indicator variable](#)

---

## What is the purpose of Data Wrangling?

Data Wrangling is the process of converting data from the initial format to a format that may be better for analysis.

## What is the fuel consumption (L/100k) rate for the diesel car?

### Import data

You can find the "Automobile Data Set" from the following link: <https://archive.ics.uci.edu/ml/machine-learning-databases/autos/imports-85.data> (<https://archive.ics.uci.edu/ml/machine-learning-databases/autos/imports-85.data>). We will be using this data set throughout this course.

### Import pandas

```
In [1]: 1 import pandas as pd
        2 import matplotlib.pyplot as plt
```

## Reading the data set from the URL and adding the related headers.

URL of the dataset

This dataset was hosted on IBM Cloud object click [HERE \(https://cocl.us/corsera\\_da0101en\\_notebook\\_bottom\)](https://cocl.us/corsera_da0101en_notebook_bottom) for free storage

```
In [2]: 1 filename = "https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/IBMDeveloperSkillsNetwork-DA0101EN-Sk
        2 <img alt="Horizontal scrollbar" data-bbox="165 625 950 645"/>
```

```
In [3]: 1 filename
```

```
Out[3]: 'https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/IBMDeveloperSkillsNetwork-DA0101EN-SkillsNetwork/labs/Data%20files/auto.csv'
```

Python list **headers** containing name of headers

```
In [4]: 1 headers = ["symboling", "normalized-losses", "make", "fuel-type", "aspiration", "num-of-doors", "body-style",
2           "drive-wheels", "engine-location", "wheel-base", "length", "width", "height", "curb-weight", "engine-type",
3           "num-of-cylinders", "engine-size", "fuel-system", "bore", "stroke", "compression-ratio", "horsepower",
4           "peak-rpm", "city-mpg", "highway-mpg", "price"]
```

Use the Pandas method **read\_csv()** to load the data from the web address. Set the parameter "names" equal to the Python list "headers".

```
In [5]: 1 df = pd.read_csv(filename, names = headers)
```

Use the method **head()** to display the first five rows of the dataframe.

```
In [6]: 1 # To see what the data set looks like, we'll use the head() method.
2 df.head()
```

Out[6]:

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	...	engine-size	fuel-system	bore	stroke	compression-ratio
0	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	2.68	9.0
1	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	2.68	9.0
2	1	?	alfa-romero	gas	std	two	hatchback	rwd	front	94.5	...	152	mpfi	2.68	3.47	9.0
3	2	164	audi	gas	std	four	sedan	fwd	front	99.8	...	109	mpfi	3.19	3.40	10.0
4	2	164	audi	gas	std	four	sedan	4wd	front	99.4	...	136	mpfi	3.19	3.40	8.0

5 rows × 26 columns



As we can see, several question marks appeared in the dataframe; those are missing values which may hinder our further analysis. So, how do we identify all those missing values and deal with them?

### How to work with missing data?

Steps for working with missing data:

1. identify missing data
2. deal with missing data
3. correct data format

## Identify and handle missing values

### Identify missing values

#### Convert "?" to NaN

In the car dataset, missing data comes with the question mark "?". We replace "?" with NaN (Not a Number), which is Python's default missing value marker, for reasons of computational speed and convenience. Here we use the function:

```
.replace(A, B, inplace = True)
```

to replace A by B

In [7]:

```
1 import numpy as np
2
3 # replace "?" to NaN
4 df.replace("?", np.nan, inplace = True)
5 df.head(5)
```

Out[7]:

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	...	engine-size	fuel-system	bore	stroke	compression-ratio
0	3	NaN	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	2.68	9.0
1	3	NaN	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	2.68	9.0
2	1	NaN	alfa-romero	gas	std	two	hatchback	rwd	front	94.5	...	152	mpfi	2.68	3.47	9.0
3	2	164	audi	gas	std	four	sedan	fwd	front	99.8	...	109	mpfi	3.19	3.40	10.0
4	2	164	audi	gas	std	four	sedan	4wd	front	99.4	...	136	mpfi	3.19	3.40	8.0

5 rows × 26 columns



Identify\_missing\_values

## Evaluating for Missing Data

The missing values are converted to default. We use the following functions to identify these missing values. There are two methods to detect missing data:

1. `.isnull()`
2. `.notnull()`

The output is a boolean value indicating whether the value that is passed into the argument is in fact missing data.

```
In [8]: 1 missing_data = df.isnull()
        2 missing_data.head(5)
```

Out[8]:

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	...	engine-size	fuel-system	bore	stroke	compression-ratio	hors
0	False	True	False	False	False	False	False	False	False	False	...	False	False	False	False	False	False
1	False	True	False	False	False	False	False	False	False	False	...	False	False	False	False	False	False
2	False	True	False	False	False	False	False	False	False	False	...	False	False	False	False	False	False
3	False	False	False	False	False	False	False	False	False	False	...	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False	False	False	...	False	False	False	False	False	False

5 rows × 26 columns



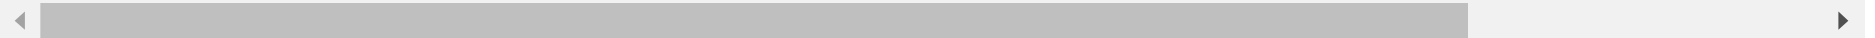
"True" stands for missing value, while "False" stands for not missing value.

```
In [9]: 1 missing_data1 = df.notnull()
        2 missing_data1.head()
```

Out[9]:

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	...	engine-size	fuel-system	bore	stroke	compression-ratio	horse
0	True	False	True	True	True	True	True	True	True	True	...	True	True	True	True	True	True
1	True	False	True	True	True	True	True	True	True	True	...	True	True	True	True	True	True
2	True	False	True	True	True	True	True	True	True	True	...	True	True	True	True	True	True
3	True	True	True	True	True	True	True	True	True	True	...	True	True	True	True	True	True
4	True	True	True	True	True	True	True	True	True	True	...	True	True	True	True	True	True

5 rows × 26 columns



"False" stands for missing value, while "True" stands for value

## Count missing values in each column

Using a for loop in Python, we can quickly figure out the number of missing values in each column. As mentioned above, "True" represents a missing value, "False" means the value is present in the dataset. In the body of the for loop the method ".value\_counts()" counts the number of "True" values.

```
In [10]: 1 for column in missing_data.columns.values.tolist():
          2     print(column)
          3     print (missing_data[column].value_counts())
          4     print("")
```

```
symboling
False      205
Name: symboling, dtype: int64
```

```
normalized-losses
False      164
True        41
Name: normalized-losses, dtype: int64
```

```
make
False      205
Name: make, dtype: int64
```

```
fuel-type
False      205
Name: fuel-type, dtype: int64
```

```
aspiration
False      205
Name: aspiration, dtype: int64
```

Based on the summary above, each column has 205 rows of data, seven columns containing missing data:

1. "normalized-losses": 41 missing data
2. "num-of-doors": 2 missing data
3. "bore": 4 missing data
4. "stroke": 4 missing data
5. "horsepower": 2 missing data

6. "peak-rpm": 2 missing data
7. "price": 4 missing data

## **Deal with missing data**

### **How to deal with missing data?**

1. drop data
  - a. drop the whole row
  - b. drop the whole column
2. replace data
  - a. replace it by mean
  - b. replace it by frequency
  - c. replace it based on other functions

Whole columns should be dropped only if most entries in the column are empty. In our dataset, none of the columns are empty enough to drop entirely. We have some freedom in choosing which method to replace data; however, some methods may seem more reasonable than others. We will apply each method to many different columns:

#### **Replace by mean:**

- "normalized-losses": 41 missing data, replace them with mean
- "stroke": 4 missing data, replace them with mean
- "bore": 4 missing data, replace them with mean
- "horsepower": 2 missing data, replace them with mean
- "peak-rpm": 2 missing data, replace them with mean

#### **Replace by frequency:**

- "num-of-doors": 2 missing data, replace them with "four".
  - Reason: 84% sedans is four doors. Since four doors is most frequent, it is most likely to occur

#### **Drop the whole row:**

- "price": 4 missing data, simply delete the whole row



- Reason: price is what we want to predict. Any data entry without price data cannot be used for prediction; therefore any row now without price data is not useful to us

### Calculate the average of the column

```
In [18]: 1 avg_loss=df["normalized-losses"].astype("float").mean(axis=0)
          2 avg_loss
          3 avr=df["normalized-losses"].replace(np.nan,avg_loss,inplace=True)
```

```
In [11]: 1 avg_norm_loss = df["normalized-losses"].astype("float").mean(axis=0)
          2 print("Average of normalized-losses:", avg_norm_loss)
```

Average of normalized-losses: 122.0

### Replace "NaN" by mean value in "normalized-losses" column

```
In [12]: 1 df["normalized-losses"].replace(np.nan, avg_norm_loss, inplace=True)
```

### Calculate the mean value for 'bore' column

```
In [13]: 1 avg_bore=df['bore'].astype('float').mean(axis=0)
          2 print("Average of bore:", avg_bore)
```

Average of bore: 3.3297512437810957

### Replace NaN by mean value

```
In [14]: 1 df["bore"].replace(np.nan, avg_bore, inplace=True)
```

## Question #1:

According to the example above, replace NaN in "stroke" column by mean.

```
In [11]: 1 # Write your code below and press Shift+Enter to execute
2 avg_stroke = df["stroke"].astype("float").mean(axis = 0)
3 print("Average of stroke", avg_stroke)
4 #replace by Nan
5 df["stroke"].replace(np.nan, avg_stroke, inplace = True)
```

Average of stroke 3.255422885572139

[Click here for the solution](#)

**Calculate the mean value for the 'horsepower' column:**

```
In [16]: 1 avg_horsepower = df['horsepower'].astype('float').mean(axis=0)
2 print("Average horsepower:", avg_horsepower)
```

Average horsepower: 104.25615763546799

**Replace "NaN" by mean value:**

```
In [17]: 1 df['horsepower'].replace(np.nan, avg_horsepower, inplace=True)
```

**Calculate the mean value for 'peak-rpm' column:**

```
In [18]: 1 avg_peakrpm=df['peak-rpm'].astype('float').mean(axis=0)
2 print("Average peak rpm:", avg_peakrpm)
```

Average peak rpm: 5125.369458128079

**Replace NaN by mean value:**

replace NaN by mean value.

```
In [19]: 1 df['peak-rpm'].replace(np.nan, avg_peakrpm, inplace=True)
```

To see which values are present in a particular column, we can use the ".value\_counts()" method:

```
In [20]: 1 df['num-of-doors'].value_counts()
```

```
Out[20]: four      114  
         two       89  
         Name: num-of-doors, dtype: int64
```

We can see that four doors are the most common type. We can also use the ".idxmax()" method to calculate for us the most common type automatically:

```
In [21]: 1 df['num-of-doors'].value_counts().idxmax()
```

```
Out[21]: 'four'
```

The replacement procedure is very similar to what we have seen previously

```
In [22]: 1 #replace the missing 'num-of-doors' values by the most frequent  
         2 df["num-of-doors"].replace(np.nan, "four", inplace=True)
```

Finally, let's drop all rows that do not have price data:

```
In [23]: 1 # simply drop whole row with NaN in "price" column  
         2 df.dropna(subset=["price"], axis=0, inplace=True)  
         3  
         4 # reset index, because we dropped two rows  
         5 df.reset_index(drop=True, inplace=True)
```

In [24]: 1 df.head()

Out[24]:

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	...	engine-size	fuel-system	bore	stroke	compression-ratio
0	3	122	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	2.68	9.0
1	3	122	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	2.68	9.0
2	1	122	alfa-romero	gas	std	two	hatchback	rwd	front	94.5	...	152	mpfi	2.68	3.47	9.0
3	2	164	audi	gas	std	four	sedan	fwd	front	99.8	...	109	mpfi	3.19	3.40	10.0
4	2	164	audi	gas	std	four	sedan	4wd	front	99.4	...	136	mpfi	3.19	3.40	8.0

5 rows × 26 columns



**Good!** Now, we obtain the dataset with no missing values.

## Correct data format

**We are almost there!**

The last step in data cleaning is checking and making sure that all data is in the correct format (int, float, text or other).

In Pandas, we use

**.dtype()** to check the data type

**.astype()** to change the data type

**Lets list the data types for each column**

```
In [25]: 1 df.dtypes
```

```
Out[25]: symboling          int64
normalized-losses  object
make              object
fuel-type         object
aspiration        object
num-of-doors      object
body-style        object
drive-wheels      object
engine-location   object
wheel-base       float64
length           float64
width            float64
height           float64
curb-weight       int64
engine-type       object
num-of-cylinders  object
engine-size       int64
fuel-system       object
bore             object
stroke           object
compression-ratio float64
horsepower        object
peak-rpm          object
city-mpg          int64
highway-mpg       int64
price            object
dtype: object
```

As we can see above, some columns are not of the correct data type. Numerical variables should have type 'float' or 'int', and variables with strings such as categories should have type 'object'. For example, 'bore' and 'stroke' variables are numerical values that describe the engines, so we should expect them to be of the type 'float' or 'int'; however, they are shown as type 'object'. We have to convert data types into a proper format for each column using the "astype()" method.

### **Convert data types to proper format**

```
In [26]: 1 df[["bore", "stroke"]] = df[["bore", "stroke"]].astype("float")
2 df[["normalized-losses"]] = df[["normalized-losses"]].astype("int")
3 df[["price"]] = df[["price"]].astype("float")
4 df[["peak-rpm"]] = df[["peak-rpm"]].astype("float")
```

**Let us list the columns after the conversion**

```
In [27]: 1 df.dtypes
```

```
Out[27]: symboling          int64
normalized-losses      int64
make                   object
fuel-type              object
aspiration             object
num-of-doors           object
body-style             object
drive-wheels           object
engine-location        object
wheel-base            float64
length                float64
width                 float64
height                float64
curb-weight            int64
engine-type            object
num-of-cylinders       object
engine-size            int64
fuel-system            object
bore                   float64
stroke                 float64
compression-ratio      float64
horsepower             object
peak-rpm               float64
city-mpg               int64
highway-mpg            int64
price                  float64
dtype: object
```

**Wonderful!**

Now, we finally obtain the cleaned dataset with no missing values and all data in its proper format.

## Data Standardization

Data is usually collected from different agencies with different formats. (Data Standardization is also a term for a particular type of data normalization, where we subtract the mean and divide by the standard deviation)

### What is Standardization?

Standardization is the process of transforming data into a common format which allows the researcher to make the meaningful comparison.

### Example

Transform mpg to L/100km:

In our dataset, the fuel consumption columns "city-mpg" and "highway-mpg" are represented by mpg (miles per gallon) unit. Assume we are developing an application in a country that accept the fuel consumption with L/100km standard

We will need to apply **data transformation** to transform mpg into L/100km?

The formula for unit conversion is

$$\text{L/100km} = 235 / \text{mpg}$$

We can do many mathematical operations directly in Pandas.

In [28]:

1 df.head()

Out[28]:

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	...	engine-size	fuel-system	bore	stroke	compression-ratio
0	3	122	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	2.68	9.0
1	3	122	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	2.68	9.0
2	1	122	alfa-romero	gas	std	two	hatchback	rwd	front	94.5	...	152	mpfi	2.68	3.47	9.0
3	2	164	audi	gas	std	four	sedan	fwd	front	99.8	...	109	mpfi	3.19	3.40	10.0
4	2	164	audi	gas	std	four	sedan	4wd	front	99.4	...	136	mpfi	3.19	3.40	8.0

5 rows × 26 columns





In [29]:

```
1 # Convert mpg to L/100km by mathematical operation (235 divided by mpg)
2 df['city-L/100km'] = 235/df["city-mpg"]
3
4 # check your transformed data
5 df.head()
```

Out[29]:

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	...	fuel-system	bore	stroke	compression-ratio	horsepower
0	3	122	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	mpfi	3.47	2.68	9.0	...
1	3	122	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	mpfi	3.47	2.68	9.0	...
2	1	122	alfa-romero	gas	std	two	hatchback	rwd	front	94.5	...	mpfi	2.68	3.47	9.0	1
3	2	164	audi	gas	std	four	sedan	fwd	front	99.8	...	mpfi	3.19	3.40	10.0	1
4	2	164	audi	gas	std	four	sedan	4wd	front	99.4	...	mpfi	3.19	3.40	8.0	1

5 rows × 27 columns



## Question #2:

According to the example above, transform mpg to L/100km in the column of "highway-mpg", and change the name of column to "highway-L/100km".

```
In [30]: 1 # Write your code below and press Shift+Enter to execute
2 df["highway-mpg"] = 235/df["highway-mpg"]
3 df.rename(columns={"highway-mpg": "highway-L/100km"}, inplace=True)
4 df.head()
```

```
Out[30]:
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	...	fuel-system	bore	stroke	compression-ratio	horsepower
0	3	122	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	mpfi	3.47	2.68	9.0	...
1	3	122	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	mpfi	3.47	2.68	9.0	...
2	1	122	alfa-romero	gas	std	two	hatchback	rwd	front	94.5	...	mpfi	2.68	3.47	9.0	1
3	2	164	audi	gas	std	four	sedan	fwd	front	99.8	...	mpfi	3.19	3.40	10.0	1
4	2	164	audi	gas	std	four	sedan	4wd	front	99.4	...	mpfi	3.19	3.40	8.0	1

5 rows × 27 columns



[Click here for the solution](#)

## Data Normalization

### Why normalization?

Normalization is the process of transforming values of several variables into a similar range. Typical normalizations include scaling the variable so the variable average is 0, scaling the variable so the variance is 1, or scaling variable so the variable values range from 0 to 1

### Example

To demonstrate normalization, let's say we want to scale the columns "length", "width" and "height"

**Target:** would like to Normalize those variables so their value ranges from 0 to 1.

**Approach:** replace original value by (original value)/(maximum value)

```
In [31]: 1 # replace (original value) by (original value)/(maximum value)
2 df['length'] = df['length']/df['length'].max()
3 df['width'] = df['width']/df['width'].max()
```

## Questiont #3:

According to the example above, normalize the column "height".

```
In [32]: 1 # Write your code below and press Shift+Enter to execute
2 df['height'] = df['height']/df['height'].max()
```

```
In [33]: 1 df[['length', 'width', 'height']].head()
```

```
Out[33]:
```

	length	width	height
0	0.811148	0.890278	0.816054
1	0.811148	0.890278	0.816054
2	0.822681	0.909722	0.876254
3	0.848630	0.919444	0.908027
4	0.848630	0.922222	0.908027

[Click here for the solution](#)

Here we can see, we've normalized "length", "width" and "height" in the range of [0,1].

## Binning

### Why binning?

Binning is a process of transforming continuous numerical variables into discrete categorical 'bins', for grouped analysis.

### Example:

In our dataset, "horsepower" is a real valued variable ranging from 48 to 288, it has 57 unique values. What if we only care about the price difference between cars with high horsepower, medium horsepower, and little horsepower (3 types)? Can we rearrange them into three 'bins' to simplify analysis?

We will use the Pandas method 'cut' to segment the 'horsepower' column into 3 bins

## Example of Binning Data In Pandas

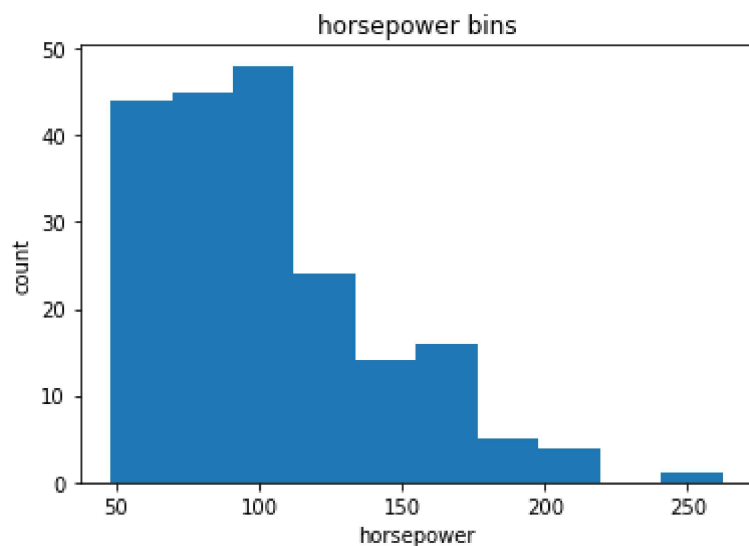
Convert data to correct format

```
In [34]: 1 df["horsepower"]=df["horsepower"].astype(int, copy=True)
```

Lets plot the histogram of horsepower, to see what the distribution of horsepower looks like.

```
In [35]: 1 %matplotlib inline
2 import matplotlib as plt
3 from matplotlib import pyplot
4 plt.pyplot.hist(df["horsepower"])
5
6 # set x/y labels and plot title
7 plt.pyplot.xlabel("horsepower")
8 plt.pyplot.ylabel("count")
9 plt.pyplot.title("horsepower bins")
```

Out[35]: Text(0.5, 1.0, 'horsepower bins')



We would like 3 bins of equal size bandwidth so we use numpy's `linspace(start_value, end_value, numbers_generated)` function.

Since we want to include the minimum value of horsepower we want to set `start_value=min(df["horsepower"])`.

Since we want to include the maximum value of horsepower we want to set `end_value=max(df["horsepower"])`.

Since we are building 3 bins of equal length, there should be 4 dividers, so `numbers_generated=4`.

We build a bin array, with a minimum value to a maximum value, with bandwidth calculated above. The bins will be values used to determine when one bin ends and another begins.

```
In [36]: 1 bins = np.linspace(min(df["horsepower"]), max(df["horsepower"]), 4)
          2 bins
```

```
Out[36]: array([ 48.          , 119.33333333, 190.66666667, 262.          ])
```

We set group names:

```
In [37]: 1 group_names = ['Low', 'Medium', 'High']
```

We apply the function "cut" to determine what each value of "df['horsepower']" belongs to.

```
In [38]: 1 df['horsepower-binned'] = pd.cut(df['horsepower'], bins, labels=group_names, include_lowest=True )
2 df[['horsepower', 'horsepower-binned']].head(20)
```

```
Out[38]:
```

	horsepower	horsepower-binned
0	111	Low
1	111	Low
2	154	Medium
3	102	Low
4	115	Low
5	110	Low
6	110	Low
7	110	Low
8	140	Medium
9	101	Low
10	101	Low
11	121	Medium
12	121	Medium
13	121	Medium
14	182	Medium
15	182	Medium
16	182	Medium
17	48	Low
18	70	Low
19	70	Low

Lets see the number of vehicles in each bin.

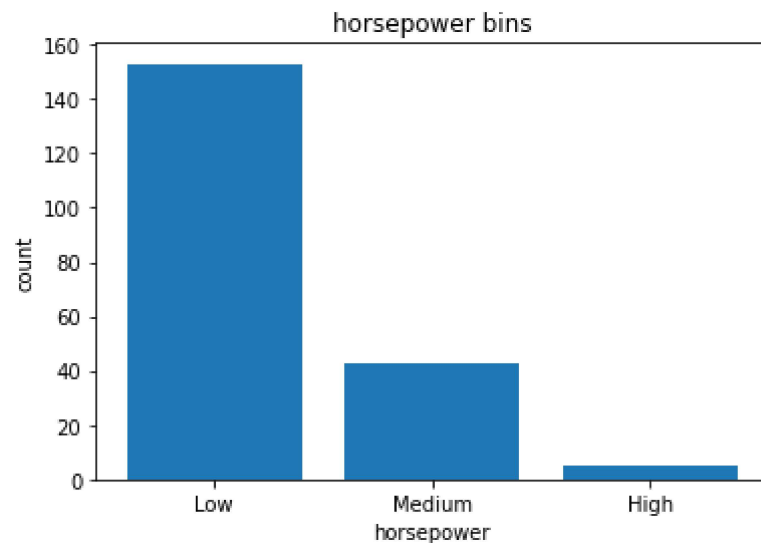
```
In [39]: 1 df["horsepower-binned"].value_counts()
```

```
Out[39]: Low      153  
Medium   43  
High      5  
Name: horsepower-binned, dtype: int64
```

Lets plot the distribution of each bin.

```
In [40]: 1 %matplotlib inline  
2 import matplotlib as plt  
3 from matplotlib import pyplot  
4 pyplot.bar(group_names, df["horsepower-binned"].value_counts())  
5  
6 # set x/y labels and plot title  
7 plt.pyplot.xlabel("horsepower")  
8 plt.pyplot.ylabel("count")  
9 plt.pyplot.title("horsepower bins")
```

```
Out[40]: Text(0.5, 1.0, 'horsepower bins')
```



Check the dataframe above carefully, you will find the last column provides the bins for "horsepower" with 3 categories ("Low","Medium" and "High").



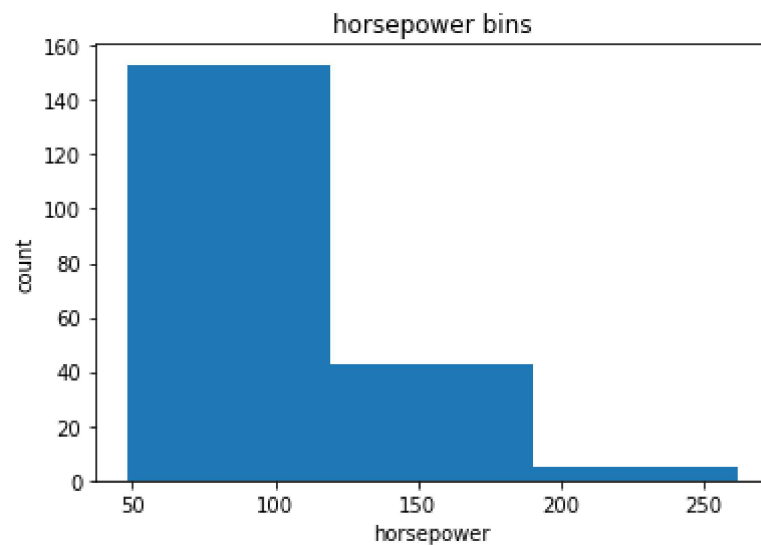
We successfully narrow the intervals from 57 to 3!

## Bins visualization

Normally, a histogram is used to visualize the distribution of bins we created above.

```
In [41]: 1 %matplotlib inline
2 import matplotlib as plt
3 from matplotlib import pyplot
4
5
6 # draw histogram of attribute "horsepower" with bins = 3
7 plt.pyplot.hist(df["horsepower"], bins = 3)
8
9 # set x/y labels and plot title
10 plt.pyplot.xlabel("horsepower")
11 plt.pyplot.ylabel("count")
12 plt.pyplot.title("horsepower bins")
```

Out[41]: Text(0.5, 1.0, 'horsepower bins')



The plot above shows the binning result for attribute "horsepower".

# Indicator variable (or dummy variable)

## What is an indicator variable?

An indicator variable (or dummy variable) is a numerical variable used to label categories. They are called 'dummies' because the numbers themselves don't have inherent meaning.

## Why we use indicator variables?

So we can use categorical variables for regression analysis in the later modules.

### Example

We see the column "fuel-type" has two unique values, "gas" or "diesel". Regression doesn't understand words, only numbers. To use this attribute in regression analysis, we convert "fuel-type" into indicator variables.

We will use the panda's method 'get\_dummies' to assign numerical values to different categories of fuel type.

```
In [42]: 1 df.columns
```

```
Out[42]: Index(['symboling', 'normalized-losses', 'make', 'fuel-type', 'aspiration',  
              'num-of-doors', 'body-style', 'drive-wheels', 'engine-location',  
              'wheel-base', 'length', 'width', 'height', 'curb-weight', 'engine-type',  
              'num-of-cylinders', 'engine-size', 'fuel-system', 'bore', 'stroke',  
              'compression-ratio', 'horsepower', 'peak-rpm', 'city-mpg',  
              'highway-mpg', 'price', 'city-L/100km', 'horsepower-binned'],  
             dtype='object')
```

get indicator variables and assign it to data frame "dummy\_variable\_1"

```
In [43]: 1 dummy_variable_1 = pd.get_dummies(df["fuel-type"])
2 dummy_variable_1.head()
```

```
Out[43]:
```

	diesel	gas
0	0	1
1	0	1
2	0	1
3	0	1
4	0	1

change column names for clarity

```
In [44]: 1 dummy_variable_1.rename(columns={'gas':'fuel-type-gas', 'diesel':'fuel-type-diesel'}, inplace=True)
2 dummy_variable_1.head()
```

```
Out[44]:
```

	fuel-type-diesel	fuel-type-gas
0	0	1
1	0	1
2	0	1
3	0	1
4	0	1

In the dataframe, column fuel-type has a value for 'gas' and 'diesel' as 0s and 1s now.

```
In [45]: 1 # merge data frame "df" and "dummy_variable_1"
2 df = pd.concat([df, dummy_variable_1], axis=1)
3
4 # drop original column "fuel-type" from "df"
5 df.drop("fuel-type", axis = 1, inplace=True)
```

In [46]: 1 df.head()

Out[46]:

	symboling	normalized-losses	make	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	length	...	compression-ratio	horsepower	peak-rpm	city-mpg
0	3	122	alfa-romero	std	two	convertible	rwd	front	88.6	0.811148	...	9.0	111	5000.0	21
1	3	122	alfa-romero	std	two	convertible	rwd	front	88.6	0.811148	...	9.0	111	5000.0	21
2	1	122	alfa-romero	std	two	hatchback	rwd	front	94.5	0.822681	...	9.0	154	5000.0	19
3	2	164	audi	std	four	sedan	fwd	front	99.8	0.848630	...	10.0	102	5500.0	24
4	2	164	audi	std	four	sedan	4wd	front	99.4	0.848630	...	8.0	115	5500.0	18

5 rows × 29 columns



The last two columns are now the indicator variable representation of the fuel-type variable. It's all 0s and 1s now.

## Question #4:

As above, create indicator variable to the column of "aspiration"

```
In [50]: 1 # Write your code below and press Shift+Enter to execute
2 dummy_variable_2 = pd.get_dummies(df['aspiration'])
3 dummy_variable_2.rename(columns={'std':'aspiration-std', 'turbo': 'aspiration-turbo'}, inplace=True)
4 dummy_variable_2.head()
```

```
Out[50]:
```

	aspiration-std	aspiration-turbo
0	1	0
1	1	0
2	1	0
3	1	0
4	1	0

[Click here for the solution](#)

## Question #5:

**Merge the new dataframe to the original dataframe then drop the column 'aspiration'**

```
In [51]: 1 # Write your code below and press Shift+Enter to execute
2 df = pd.concat([df,dummy_variable_2],axis = 1)
3 # drop original column "aspiration" from "df"
4 df.drop("aspiration",axis = 1, inplace=True)
```

[Click here for the solution](#)

Save the new csv

```
In [52]: 1 df.to_csv('clean_df.csv')
```

Type *Markdown* and LaTeX:  $\alpha^2$